

Đồ thị: Đường đi ngắn nhất

Dijkstra's algorithm

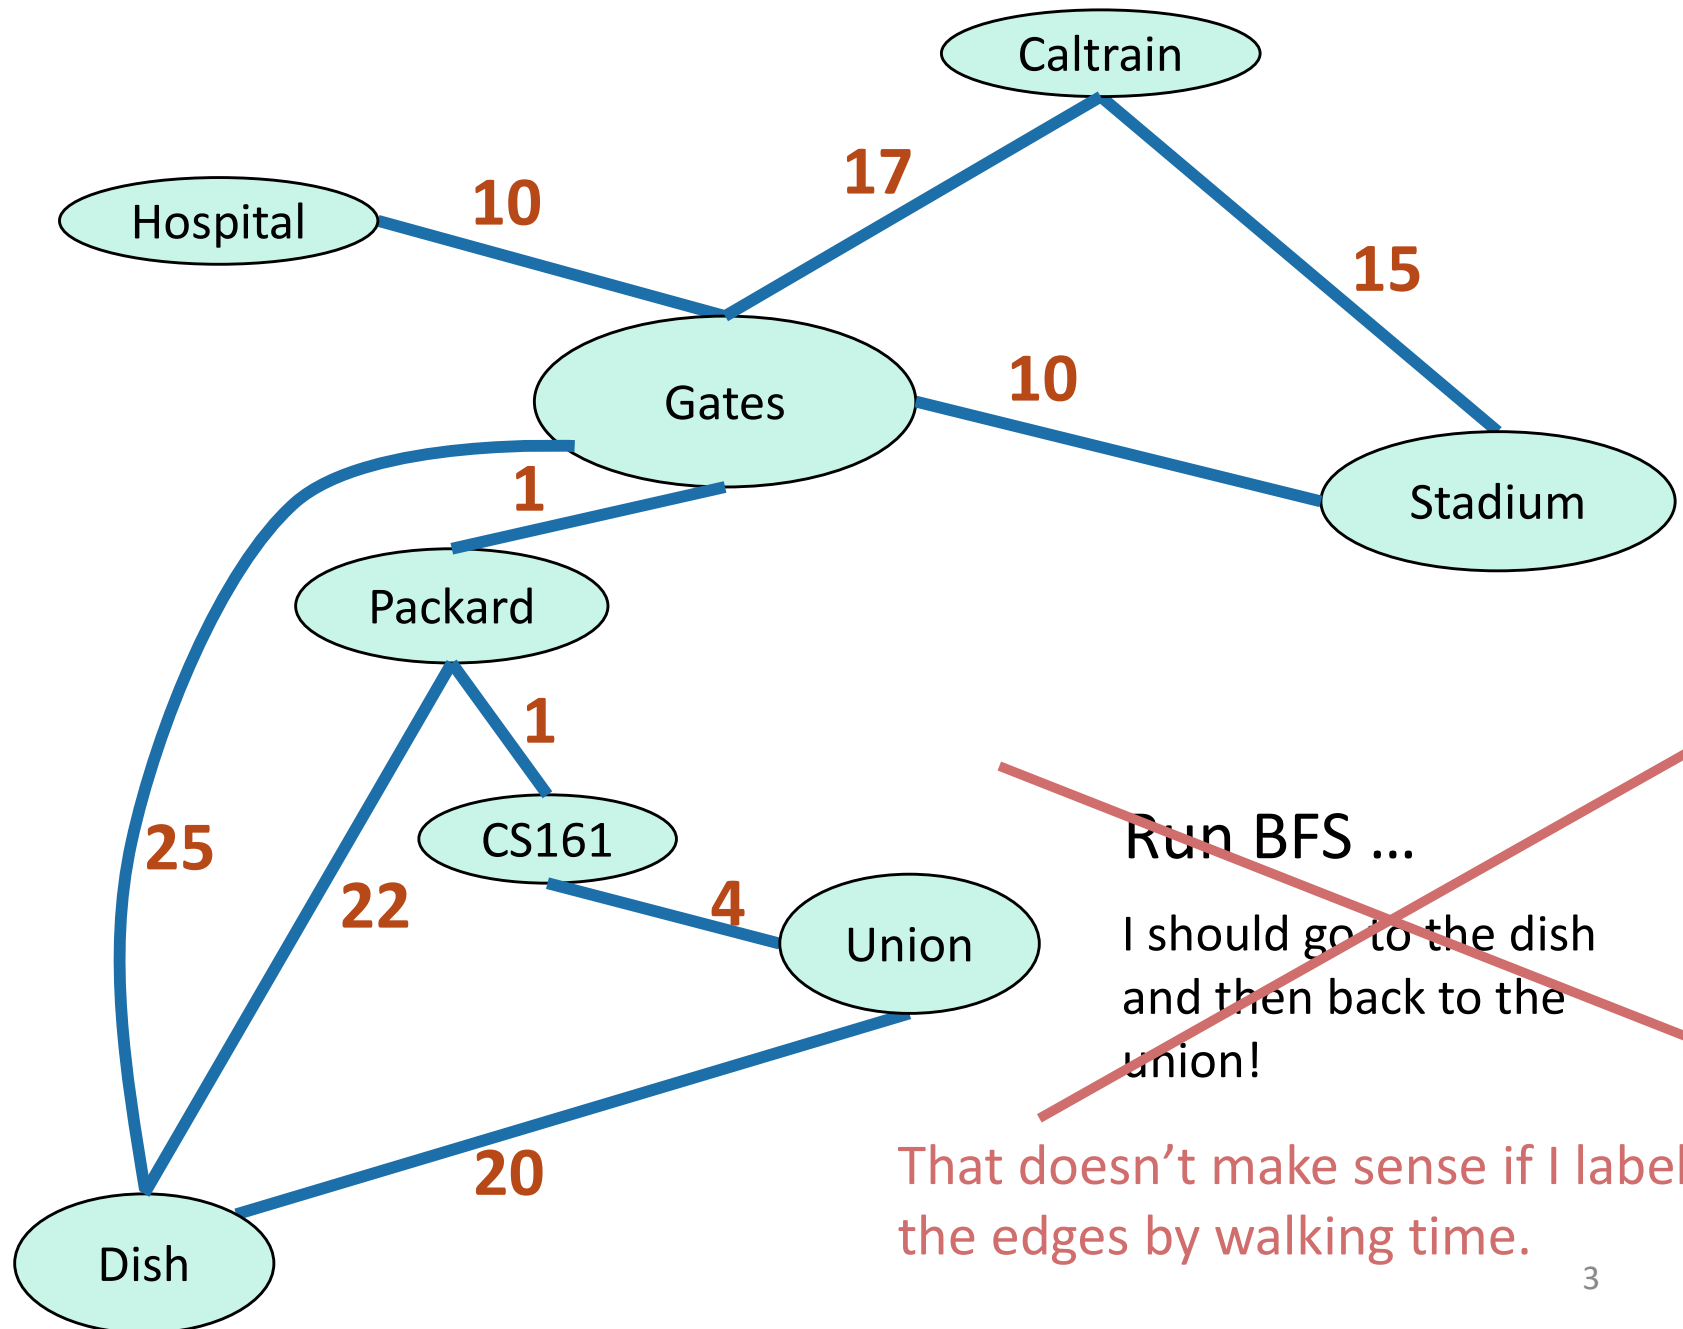


Nội dung

- Duyệt đồ thị có trọng số
- Thuật toán tìm đường đi ngắn nhất - Dijkstra's algorithm

Sử dụng một phần tài liệu bài giảng CS161 Stanford University

Shortest path from Gates to the Union?

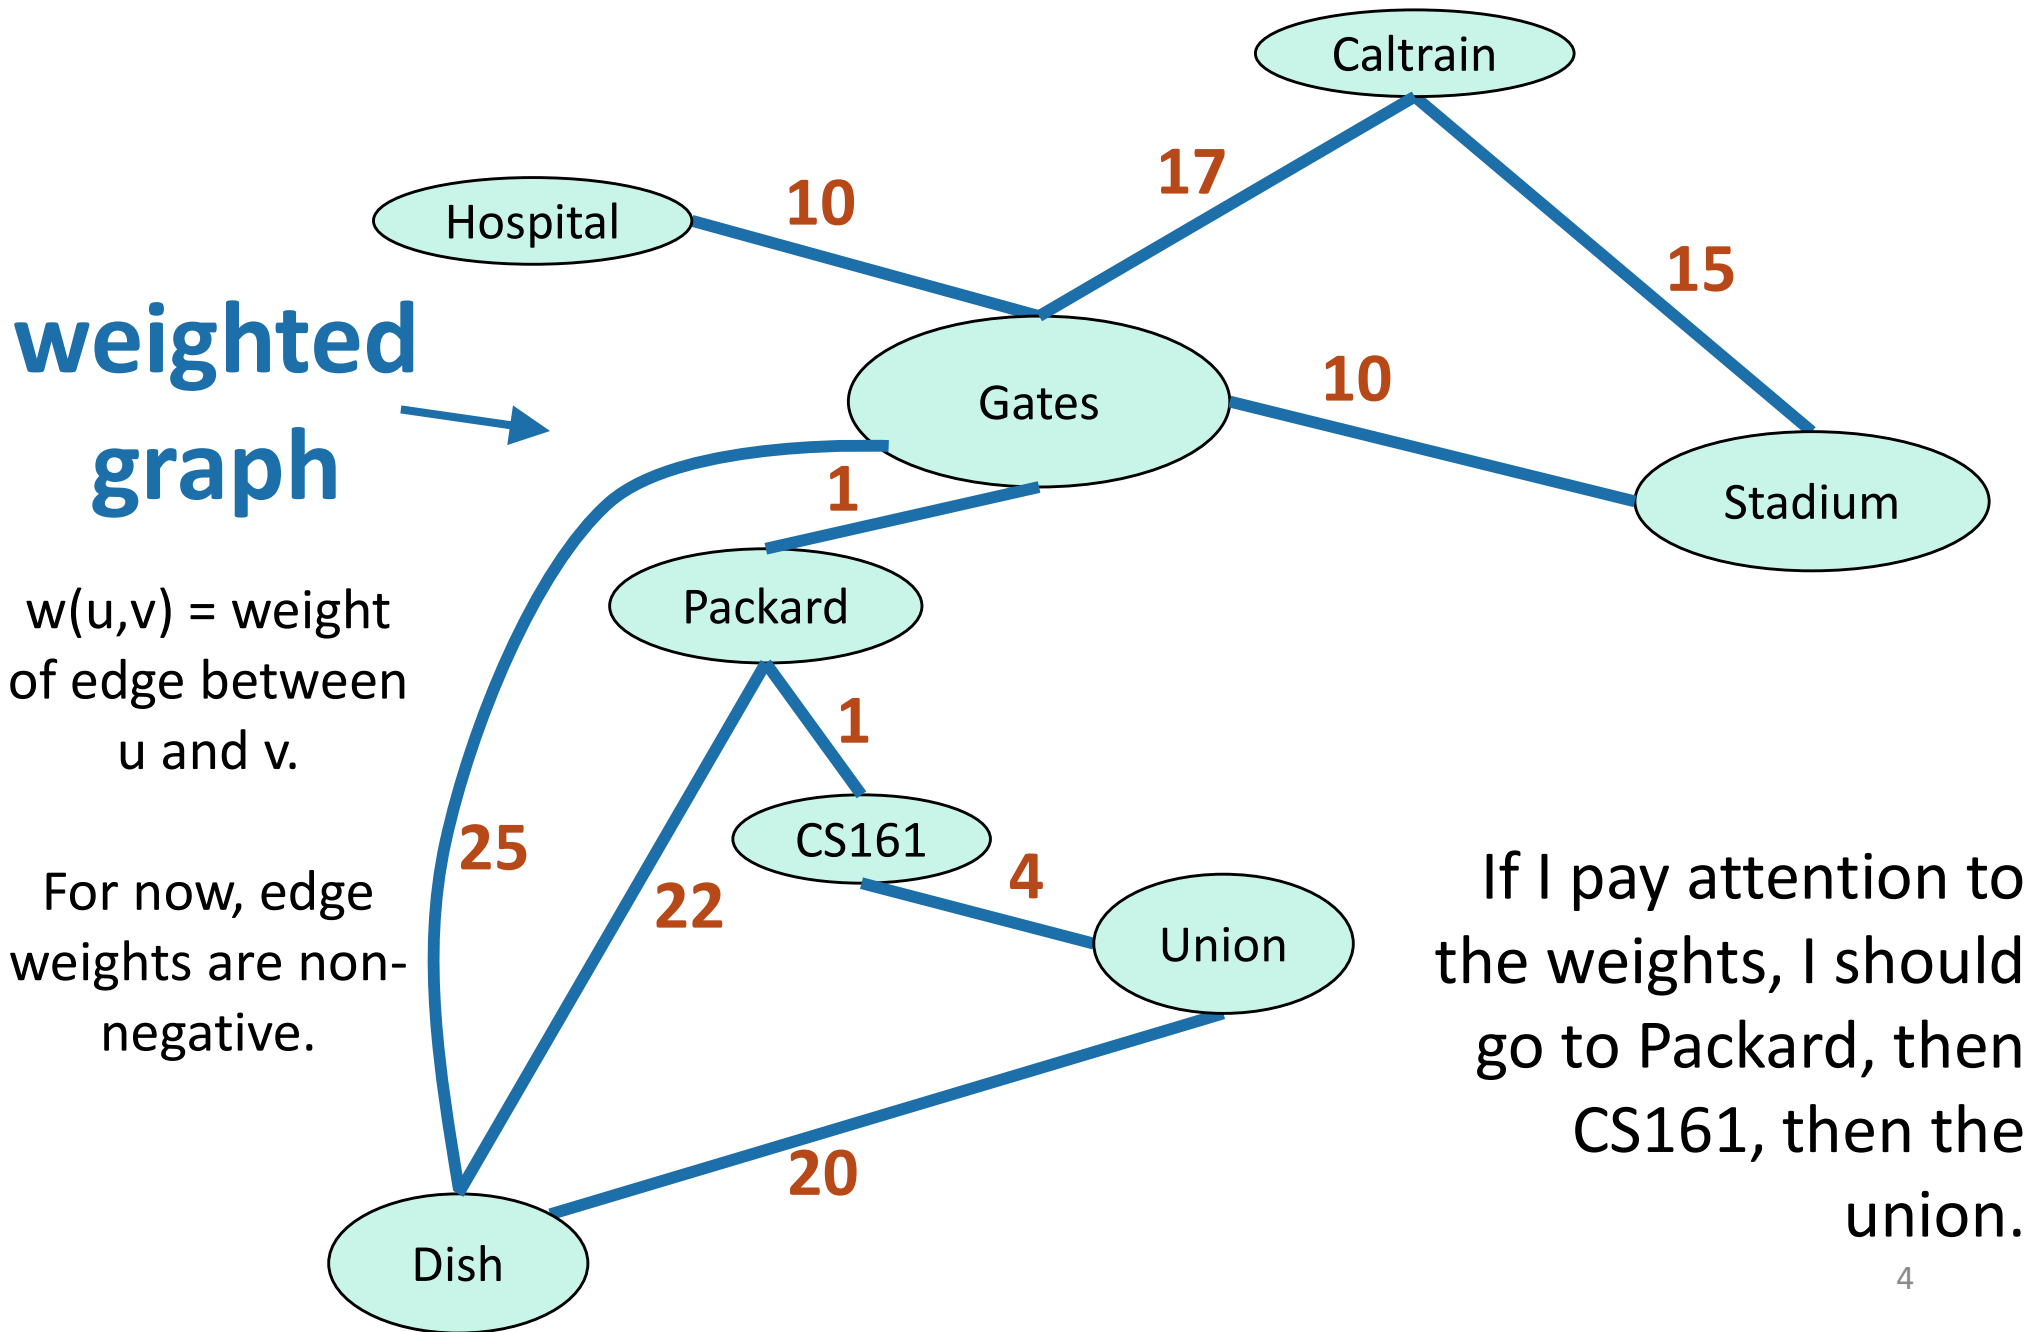


~~Run BFS ...~~

~~I should go to the dish
and then back to the
union!~~


That doesn't make sense if I label
the edges by walking time.

Shortest path from Gates to the Union?



Shortest path problem

- What is the **shortest path** between u and v in a weighted graph?
 - the **cost** of a path is the sum of the weights along that path
 - The **shortest path** is the one with the minimum cost.



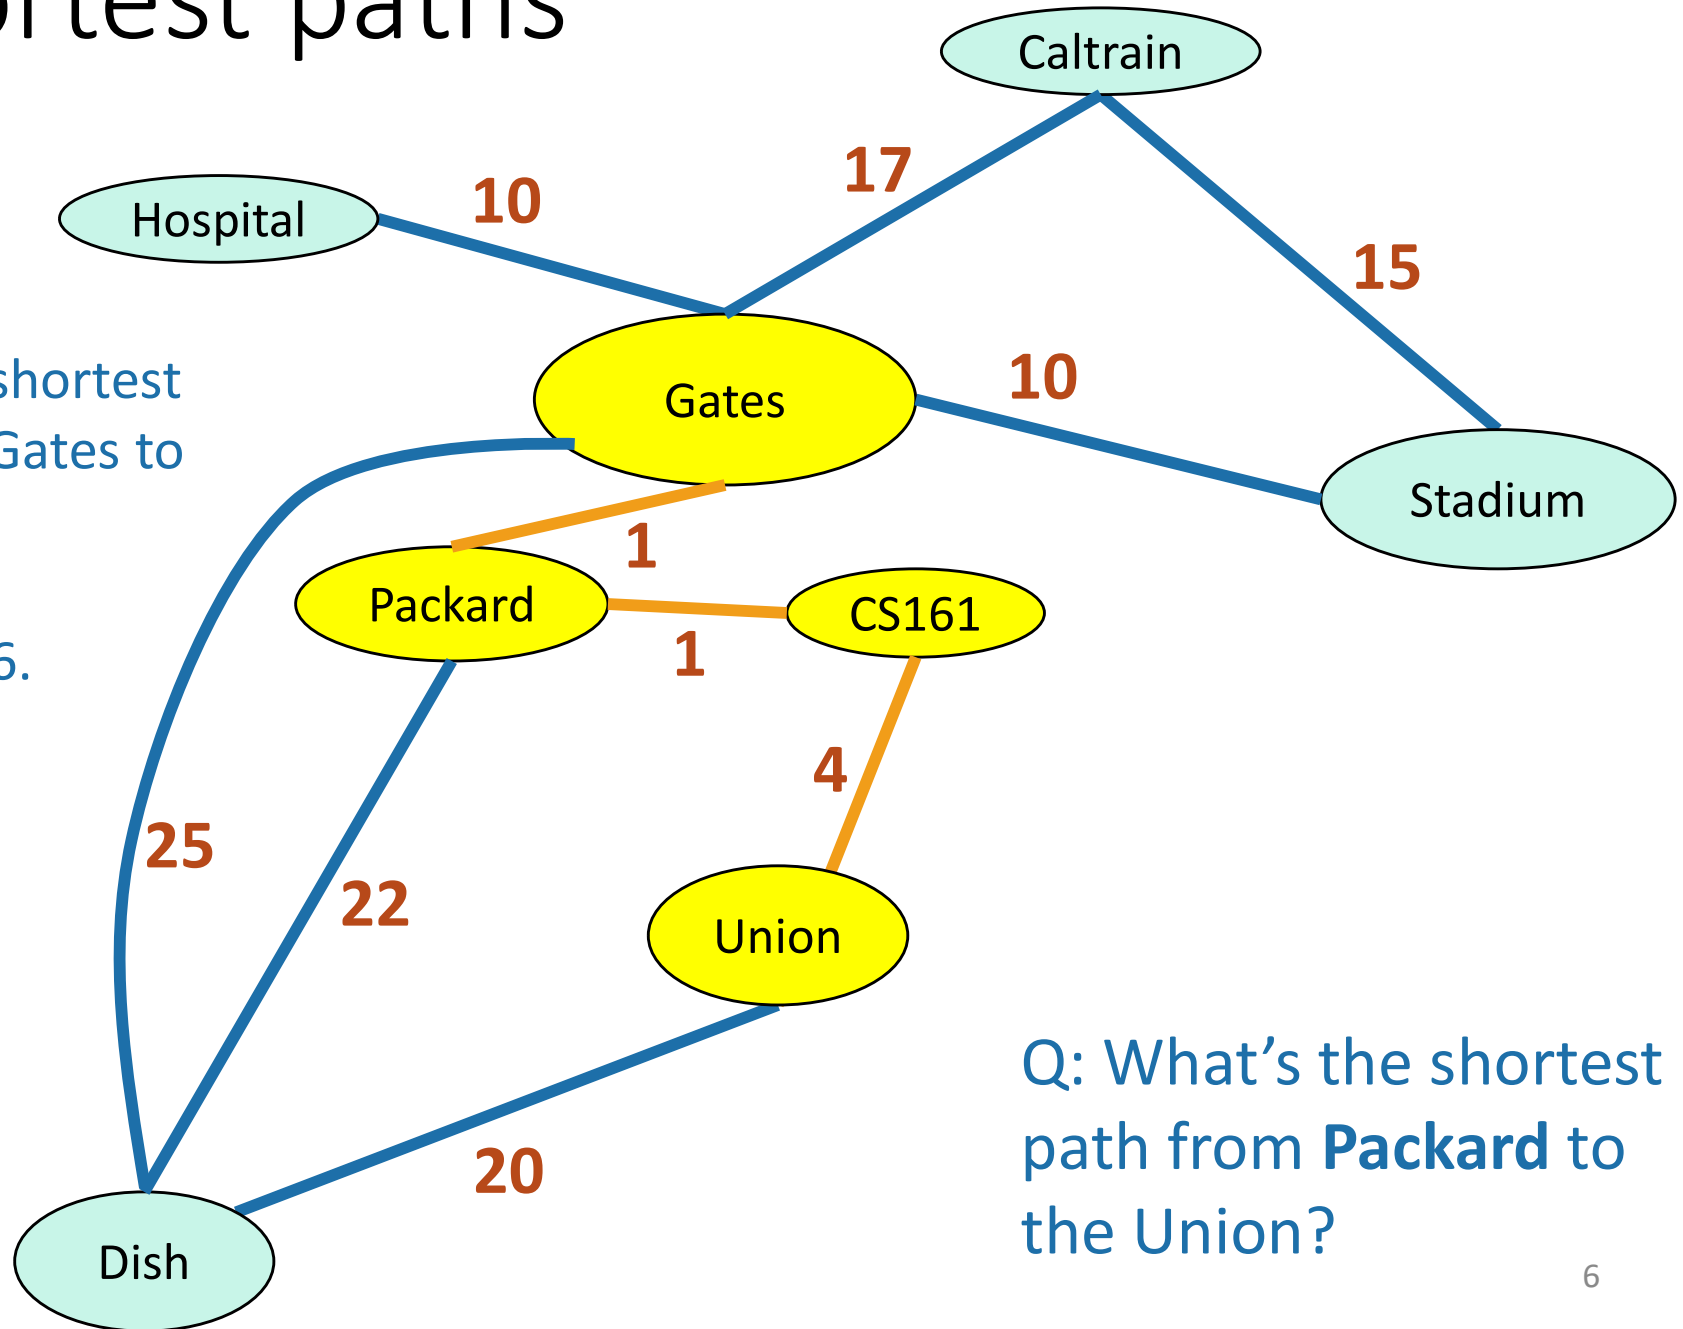
- The **distance** $d(u,v)$ between two vertices u and v is the cost of the the shortest path between u and v .
- For this lecture **all graphs are directed**, but to save on notation I'm just going to draw undirected edges.



Shortest paths

This is the shortest path from Gates to the Union.

It has cost 6.




Q: What's the shortest path from **Packard** to the Union?

Warm-up

- A sub-path of a shortest path is also a shortest path.


- Say **this** is a shortest path from s to t .
- Claim: **this** is a shortest path from s to x .
 - Suppose not, **this** one is a shorter path from s to x .
 - But then that gives an **even shorter path** from s to t !



CONTRADICTION!!

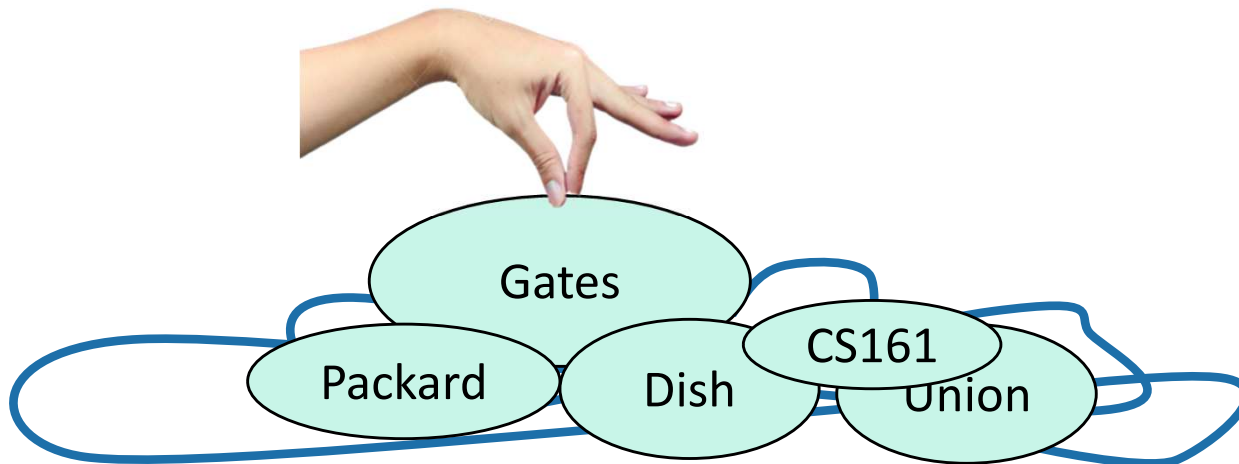
Dijkstra's algorithm

- Finds shortest paths from Gates to everywhere else.




Dijkstra

intuition




Dijkstra intuition


A vertex is done when it's not on the ground anymore.



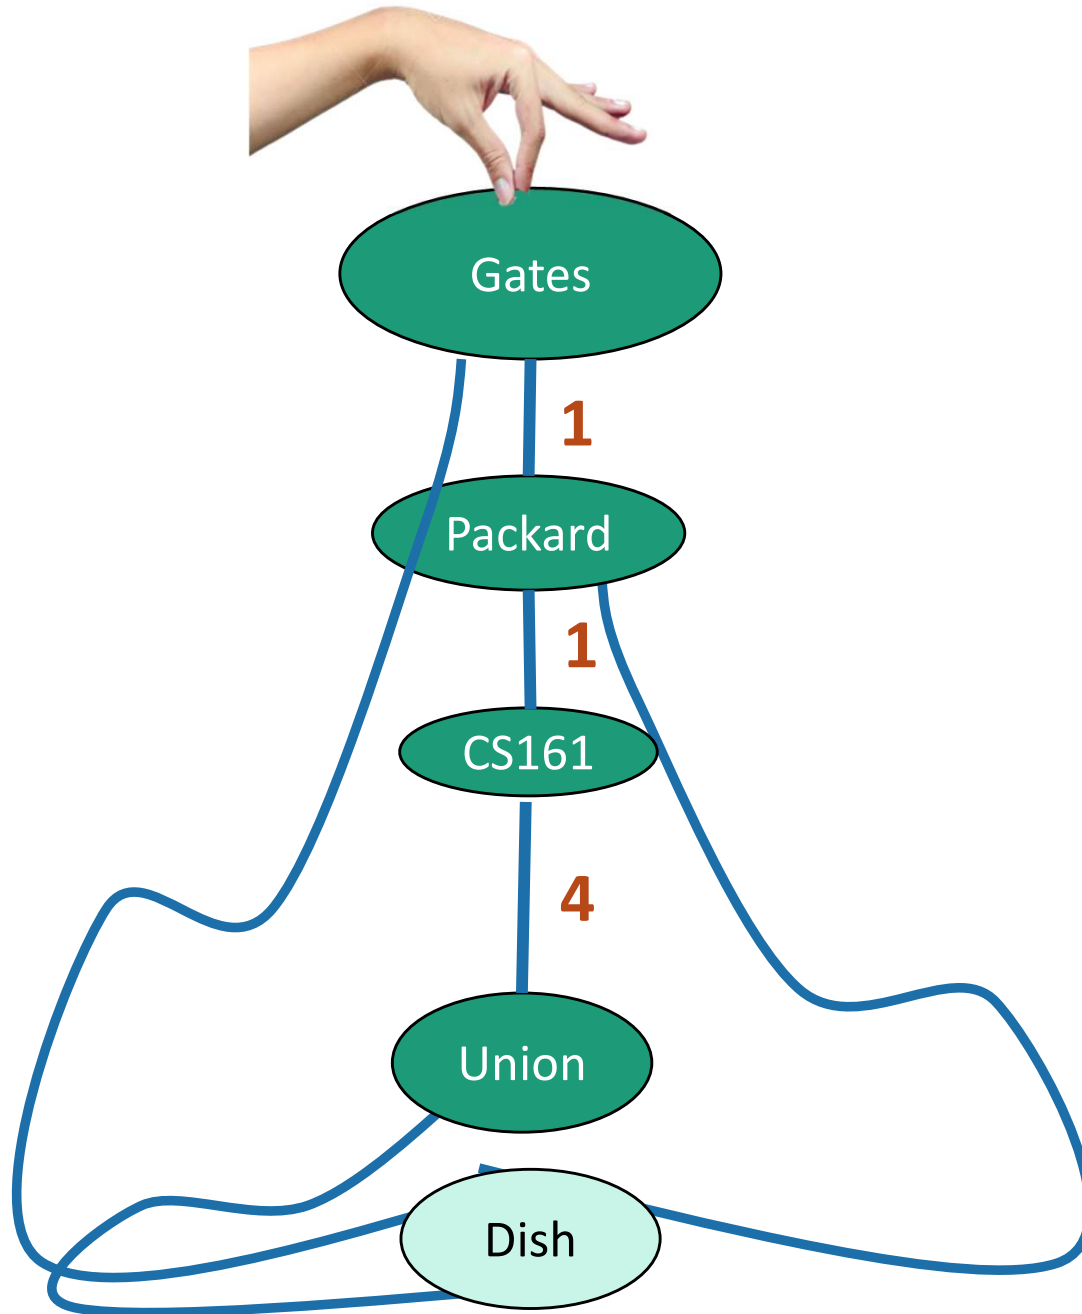
Dijkstra intuition



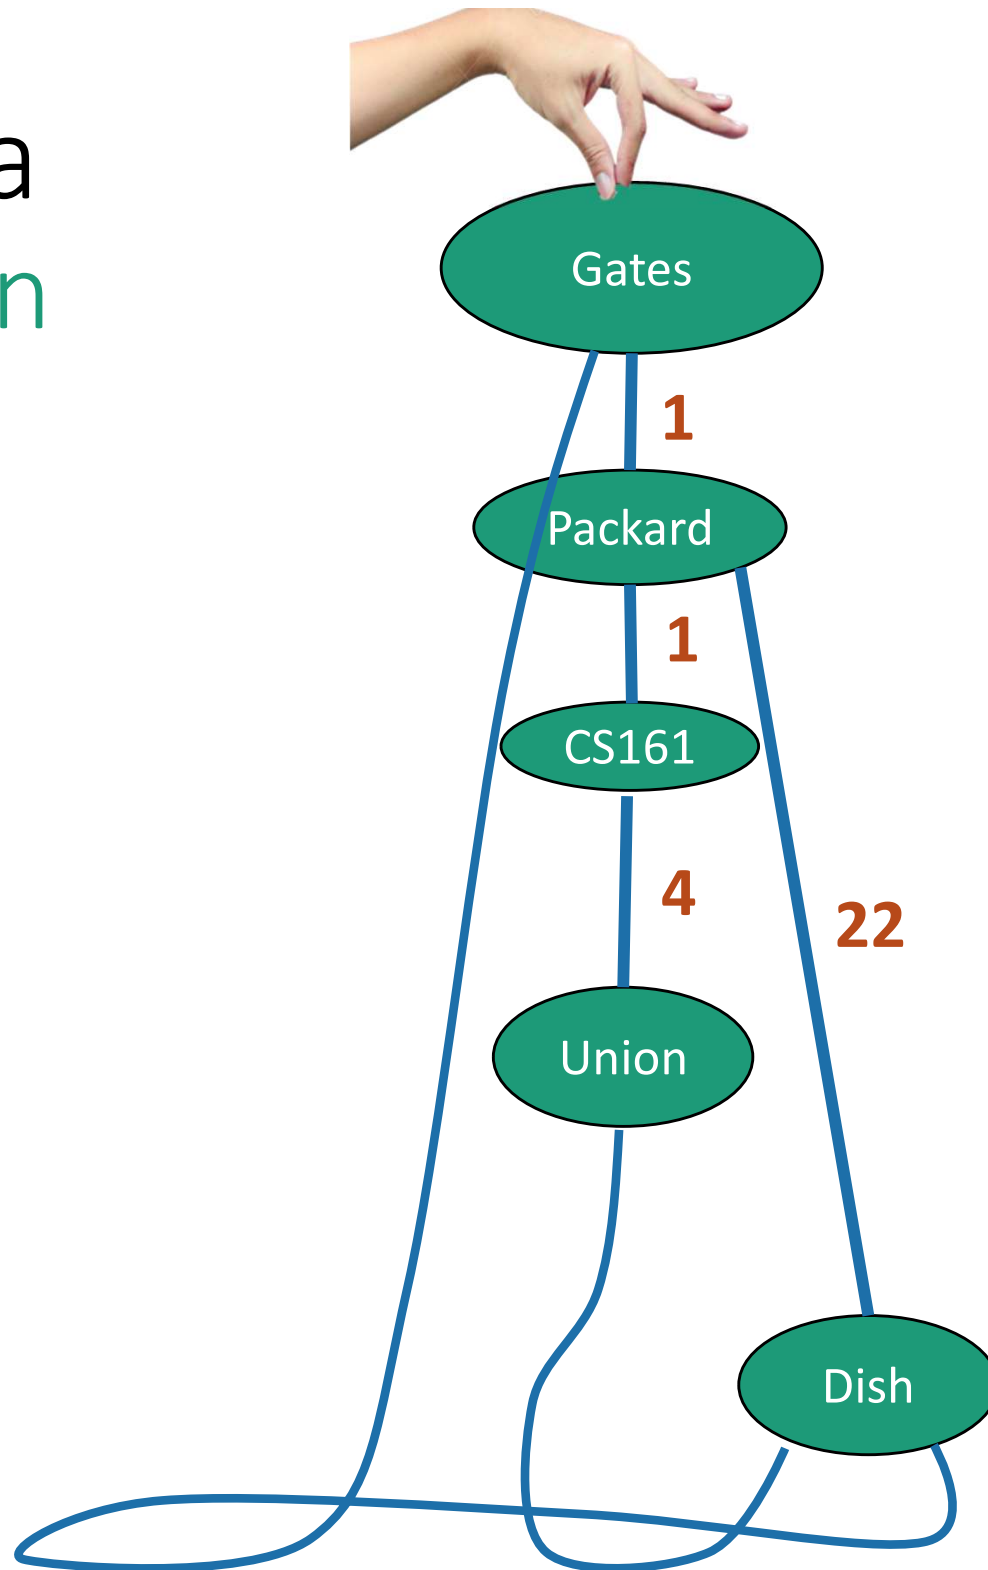
Dijkstra intuition



Dijkstra intuition



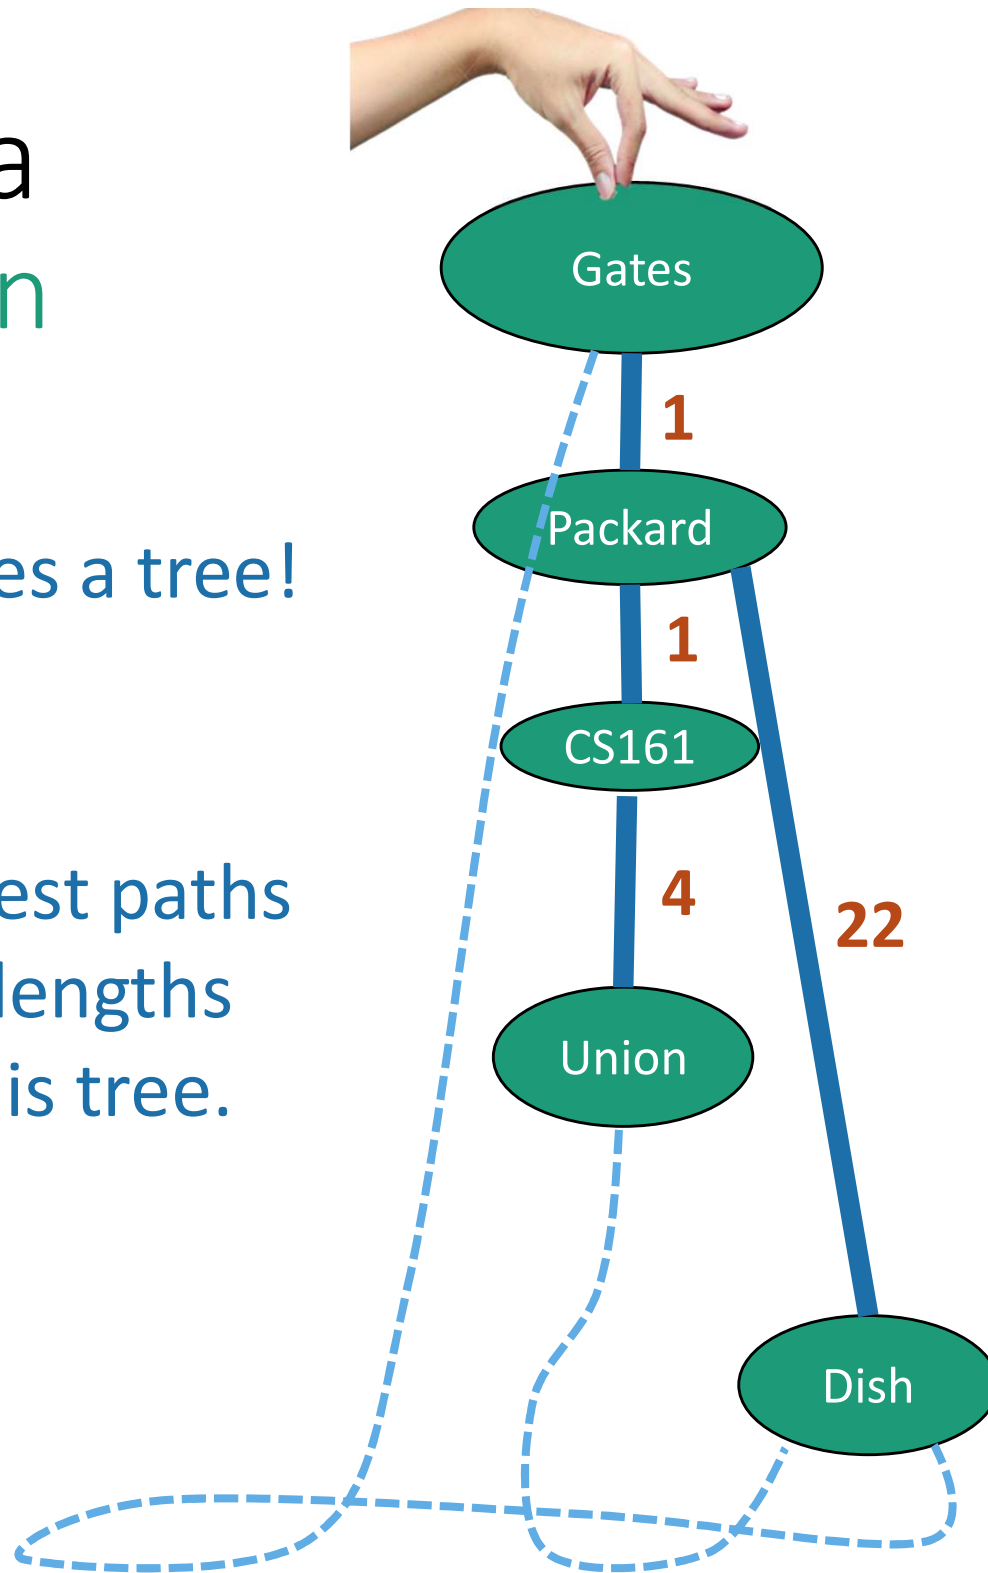
Dijkstra intuition



Dijkstra intuition


This creates a tree!

The shortest paths
are the lengths
along this tree.




How do we actually implement this?

- **Without** string and gravity?




Dijkstra by example

How far is a node from Gates?



I'm not sure yet



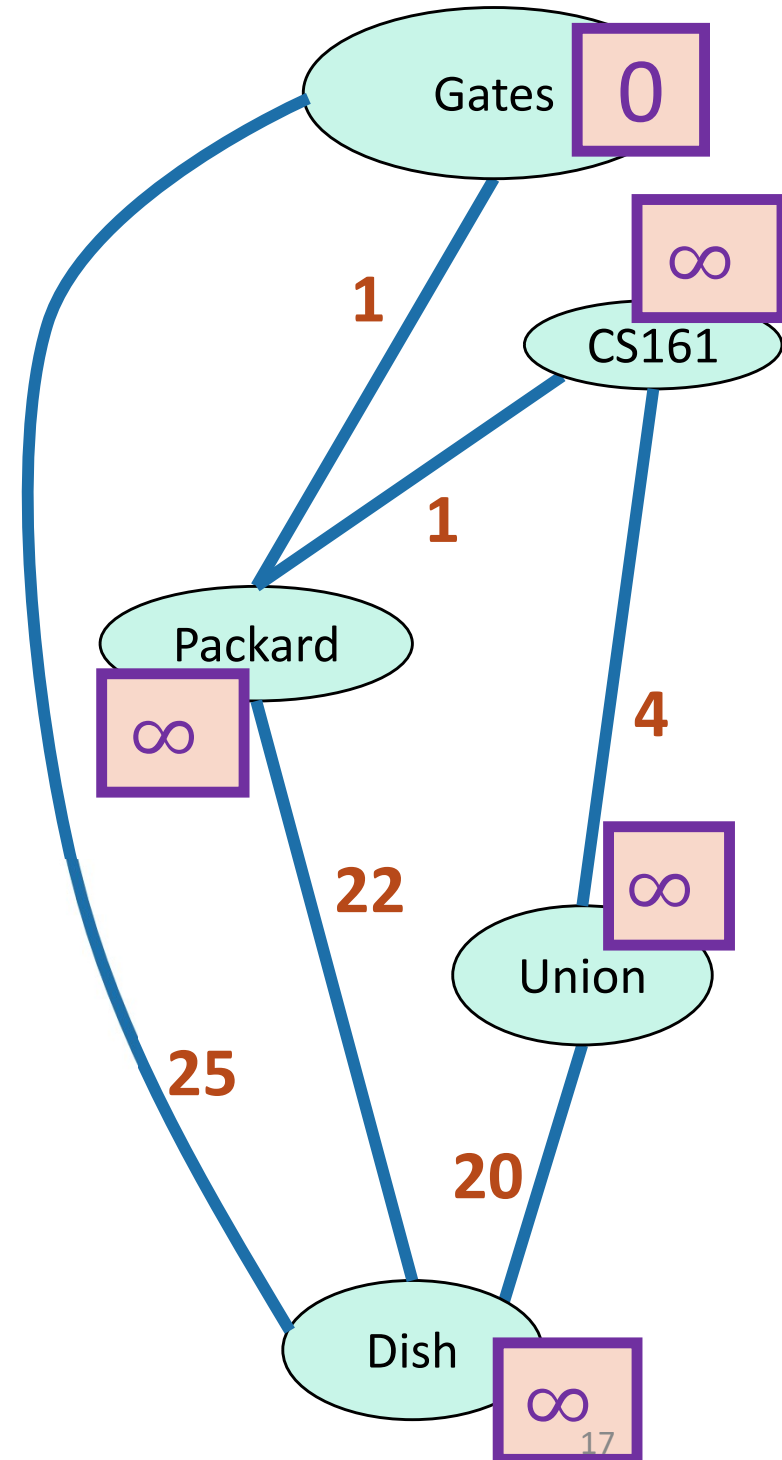
I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.


Initialize $d[v] = \infty$
for all non-starting vertices v ,
and $d[\text{Gates}] = 0$

- Pick the **not-sure** node u with the smallest estimate $d[u]$.




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

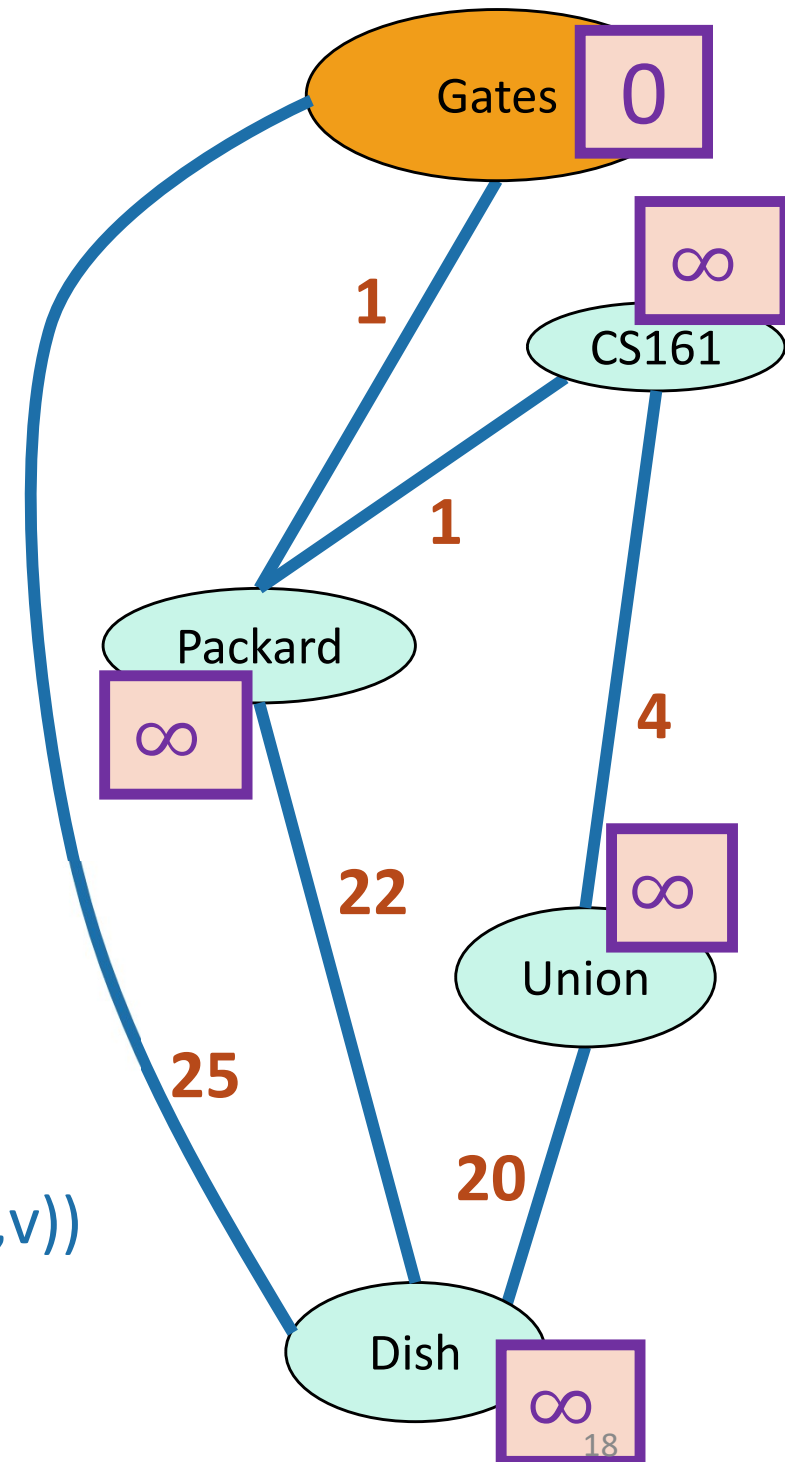


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

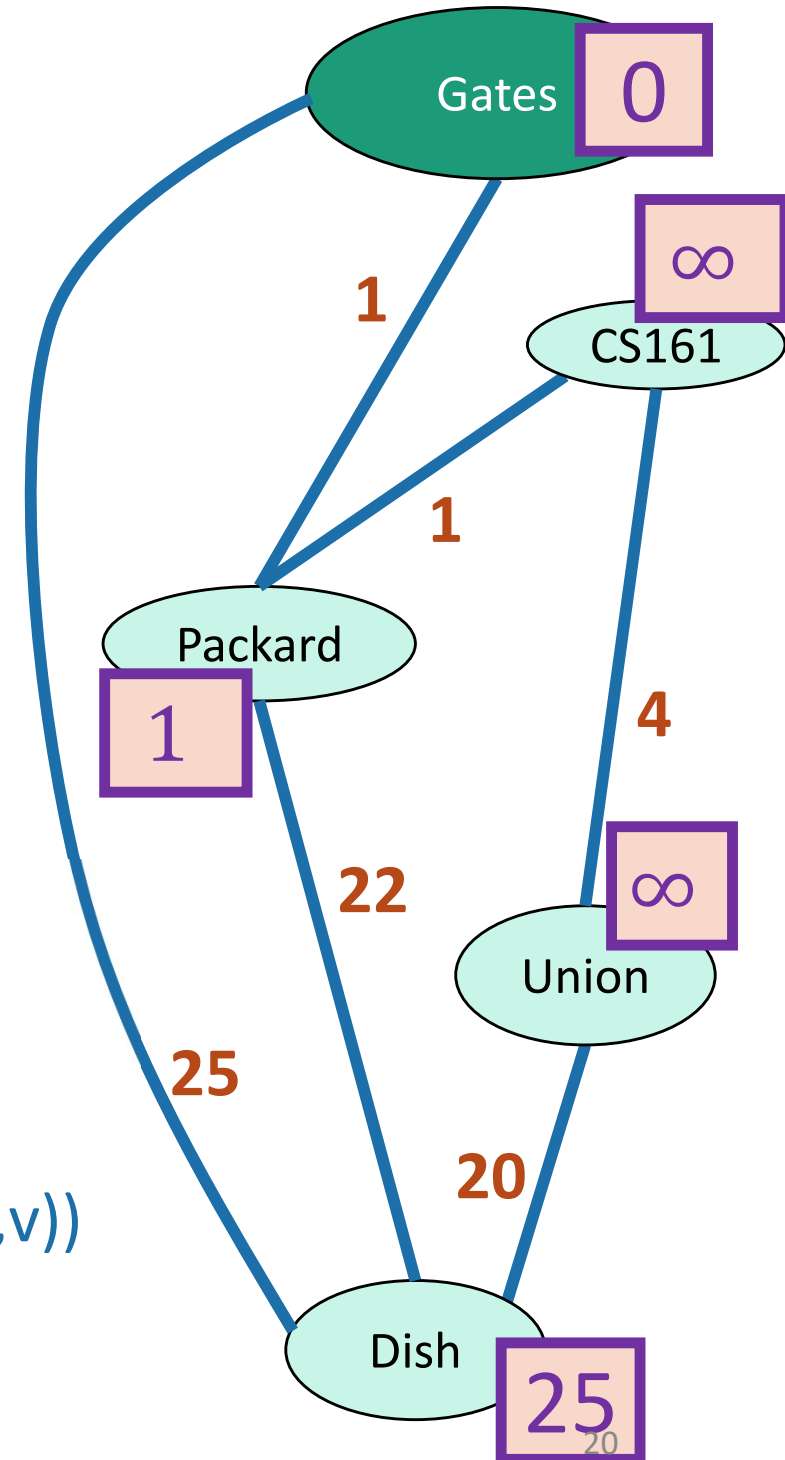


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure



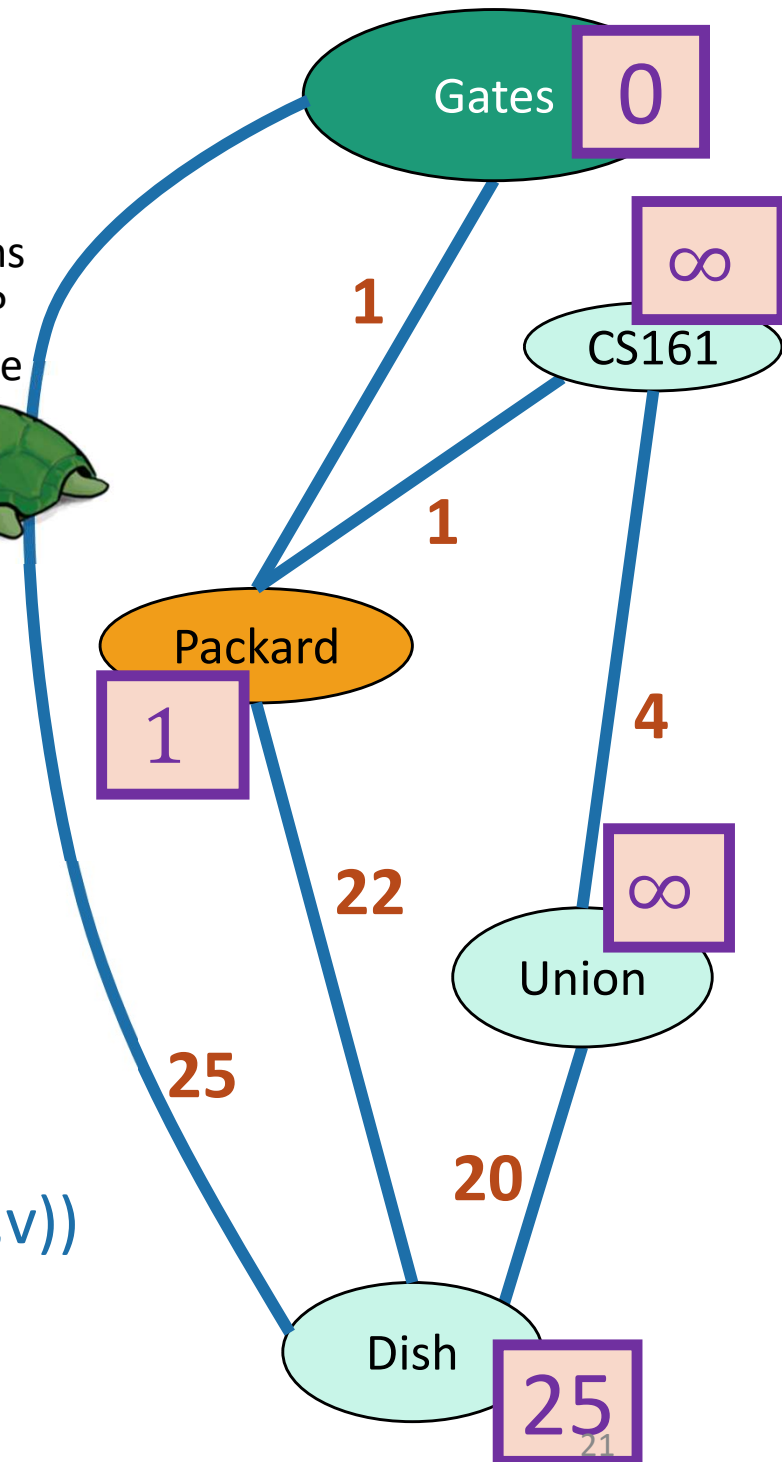

$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



Current node u


- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat

Packard has three neighbors. What happens when we update them?
1 min. think; 1 min. share




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure



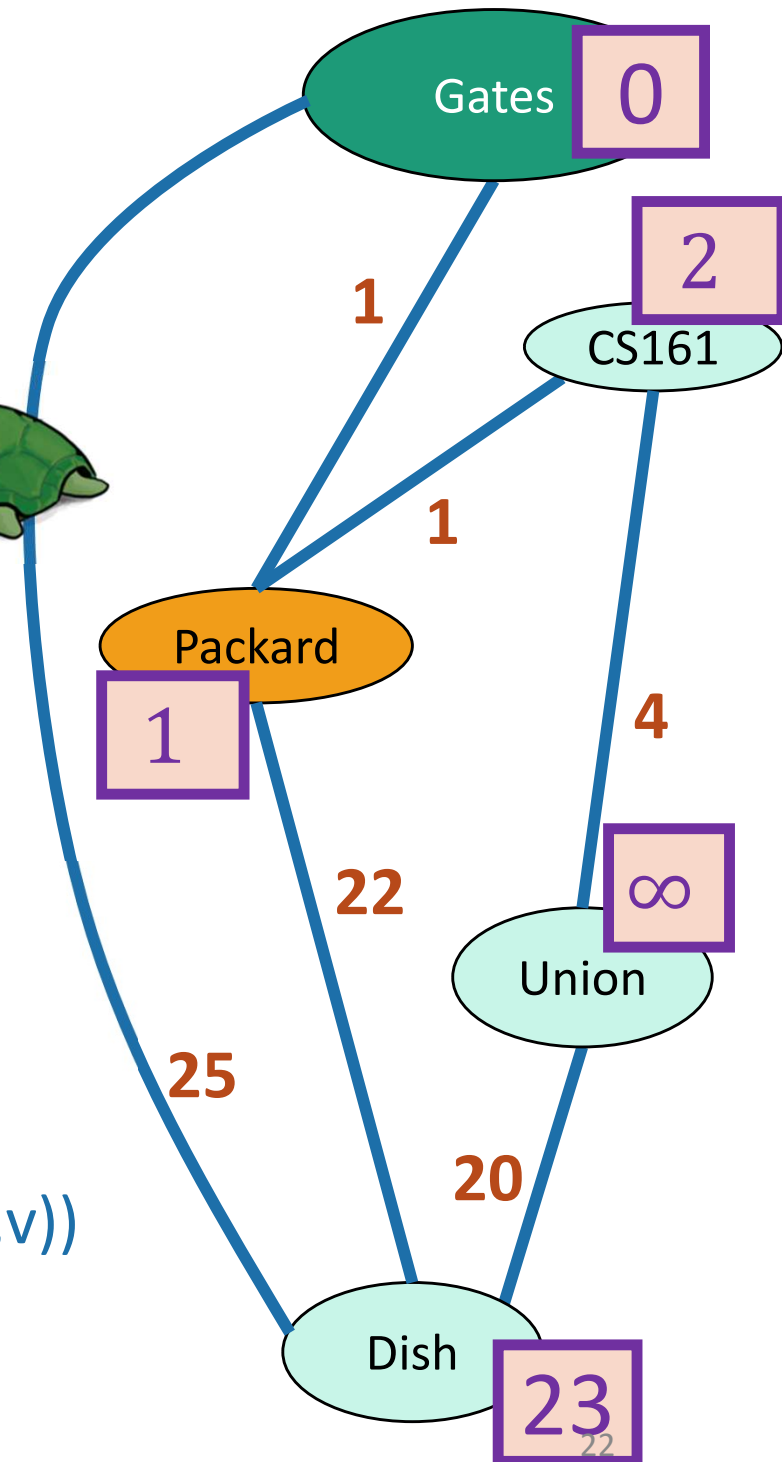

$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



Current node u


- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat

Packard has three neighbors. What happens when we update them?




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

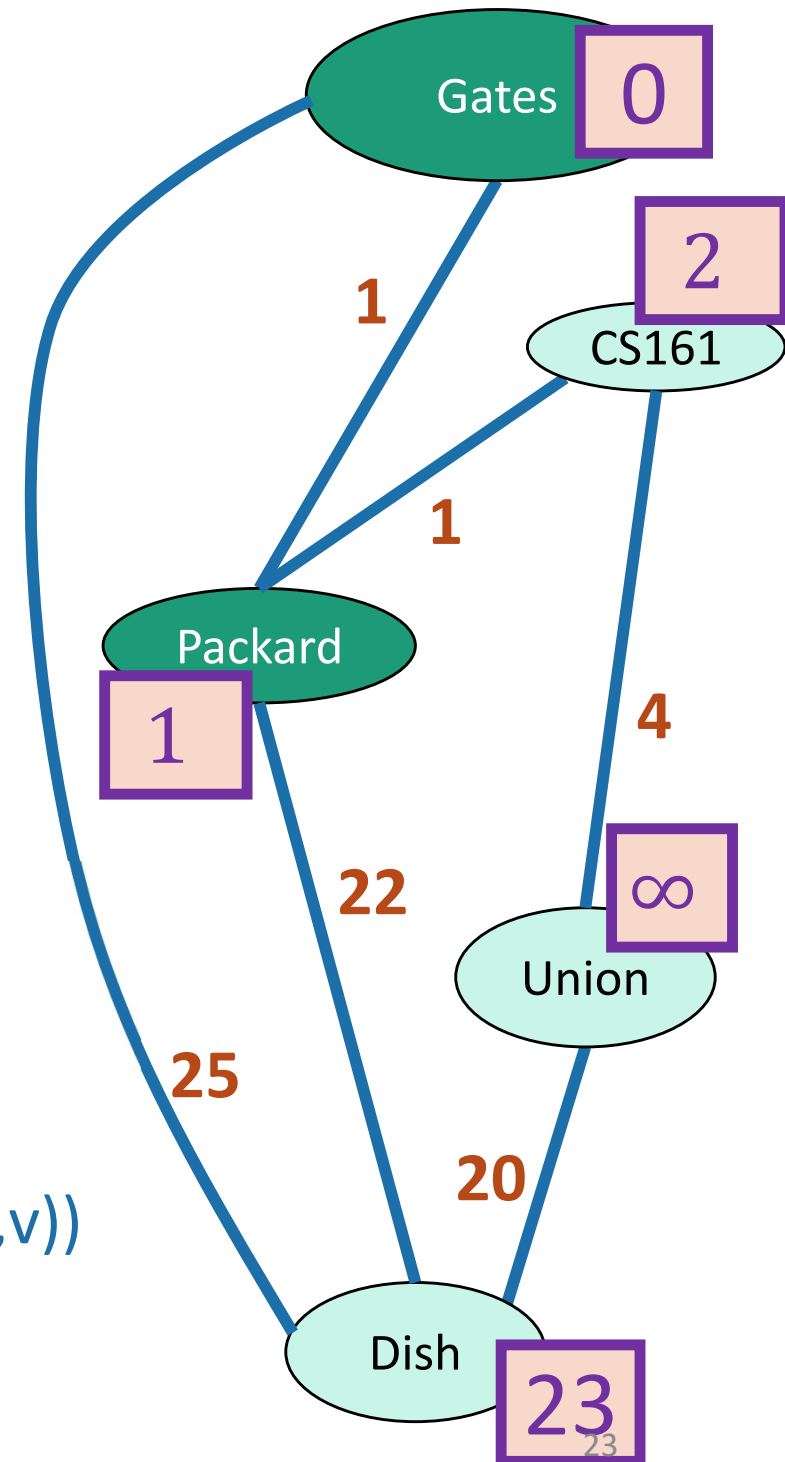


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

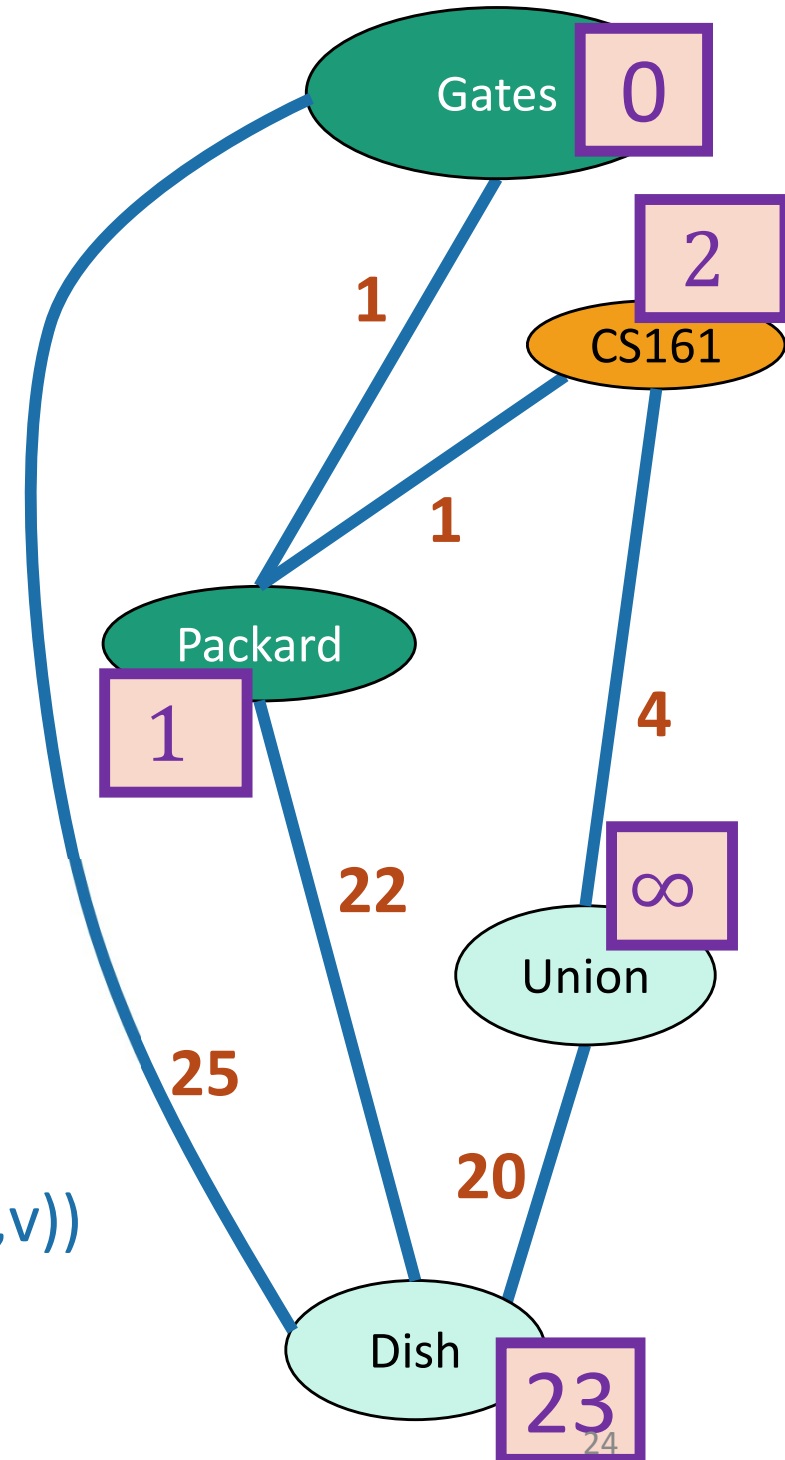


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

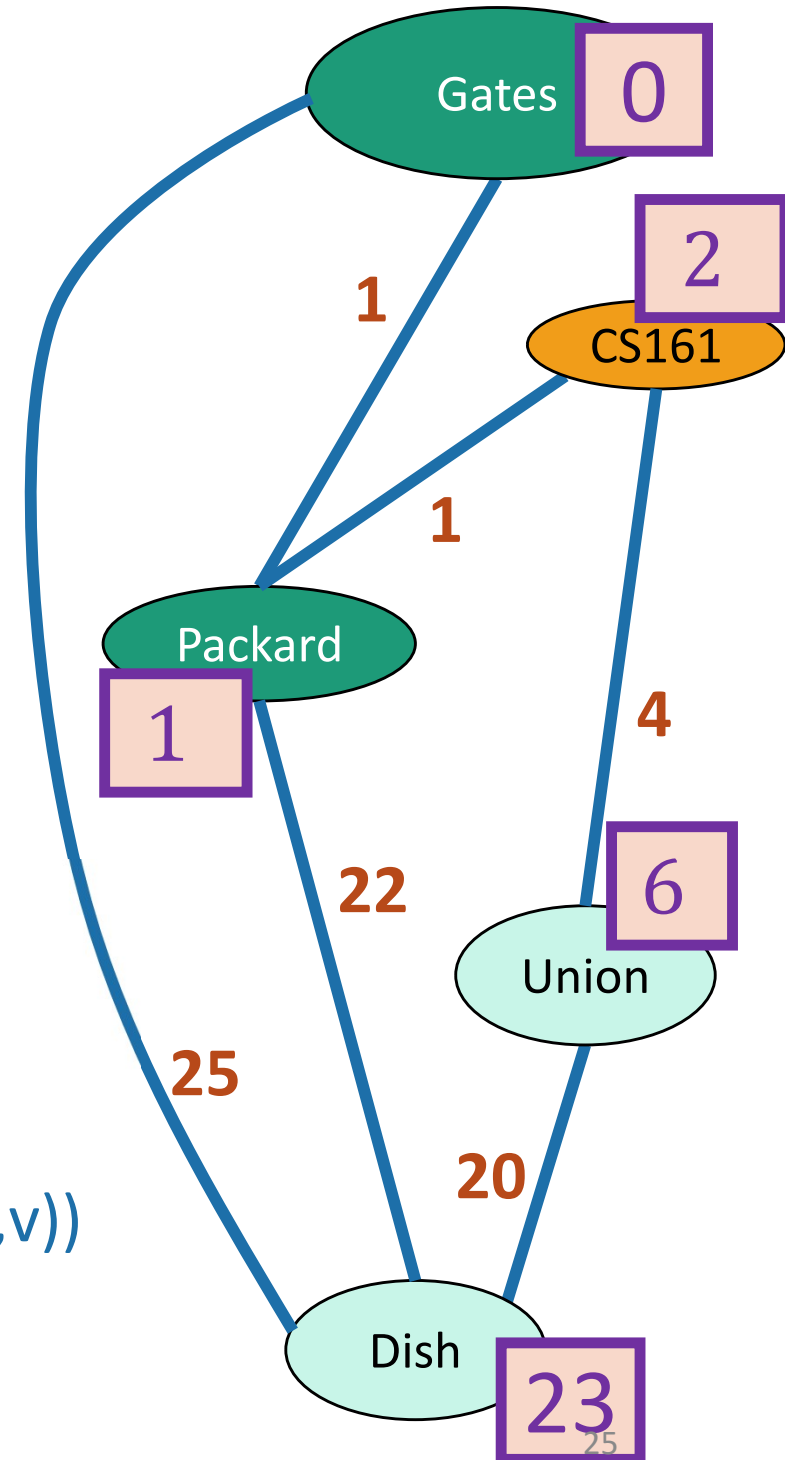


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

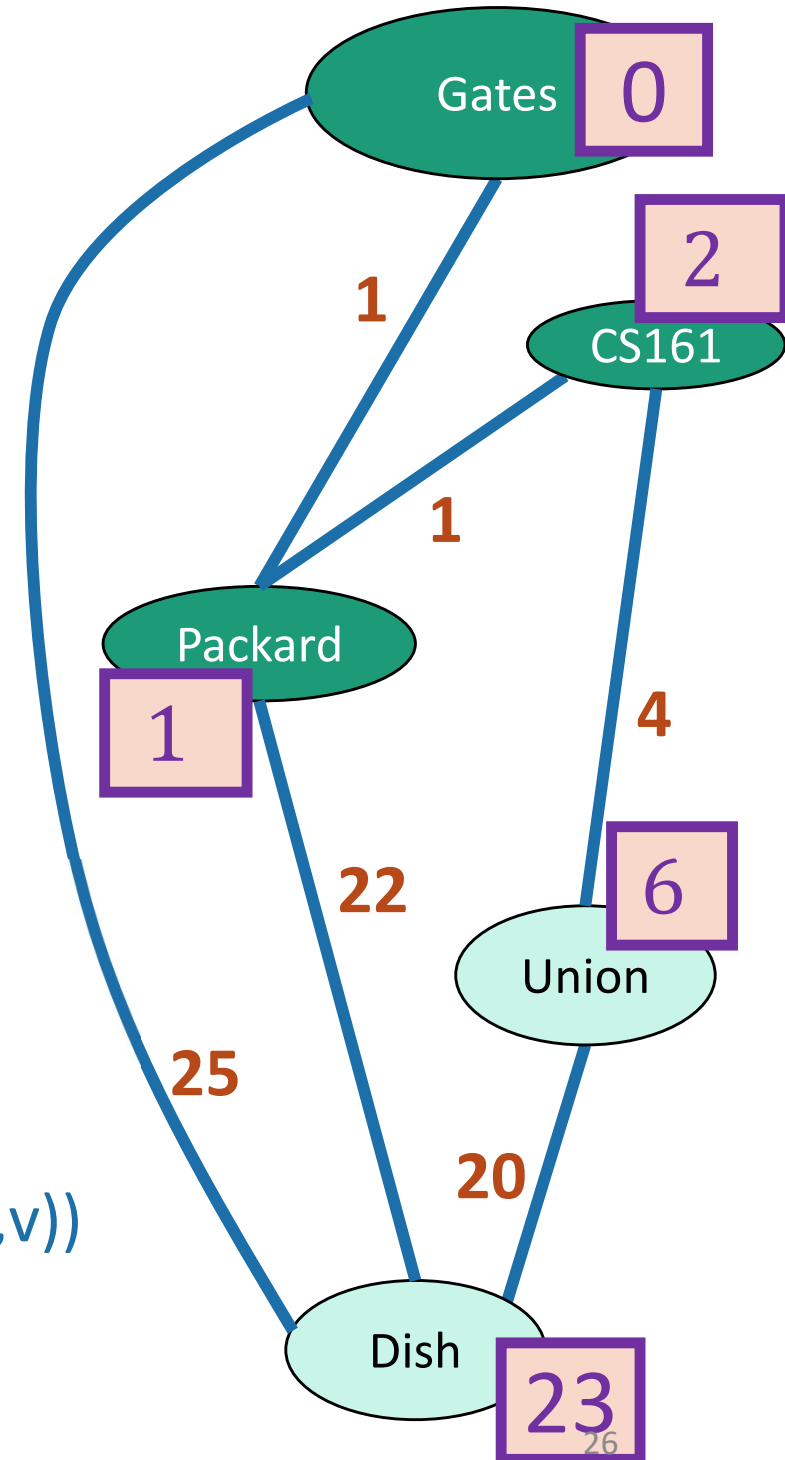


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

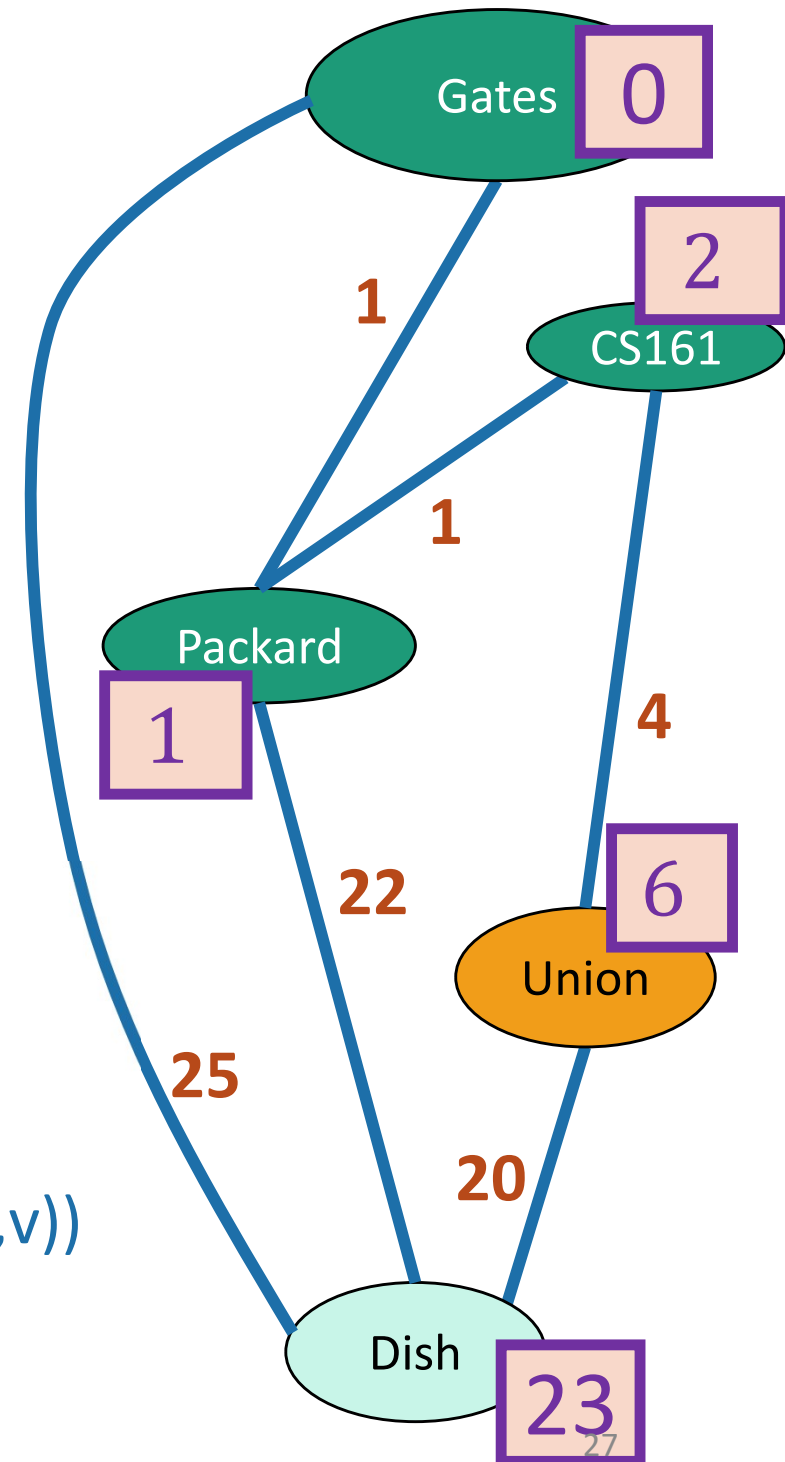


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

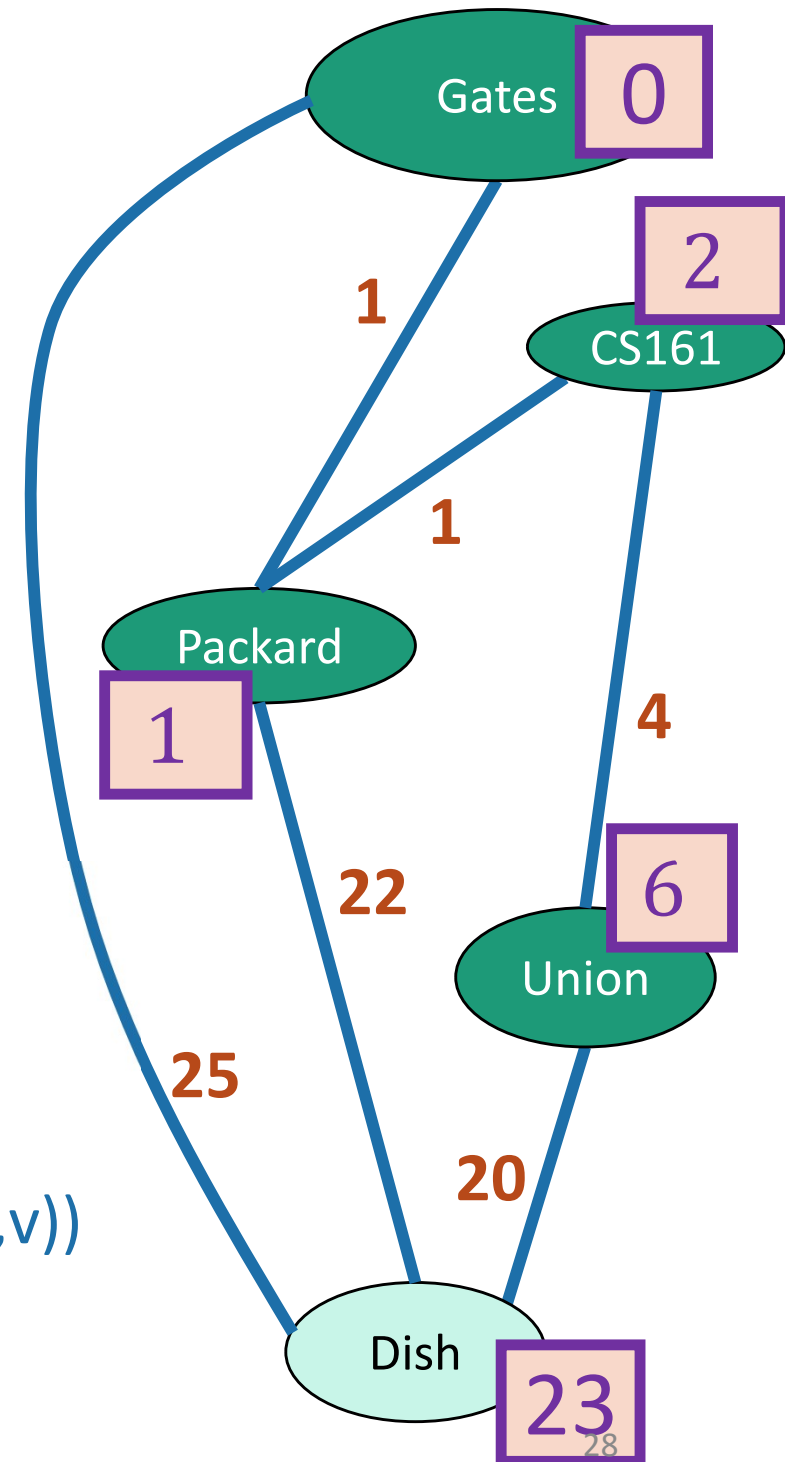


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

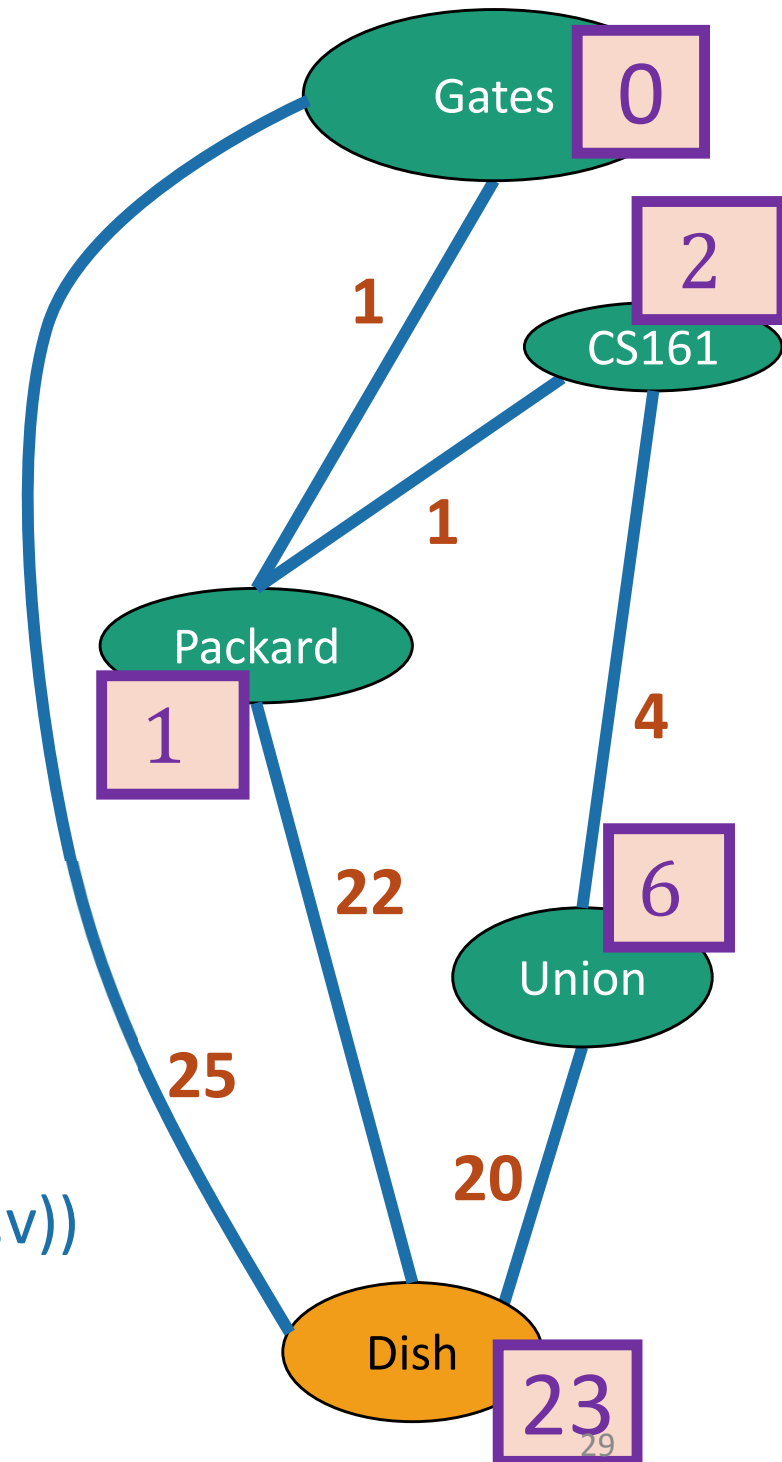


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

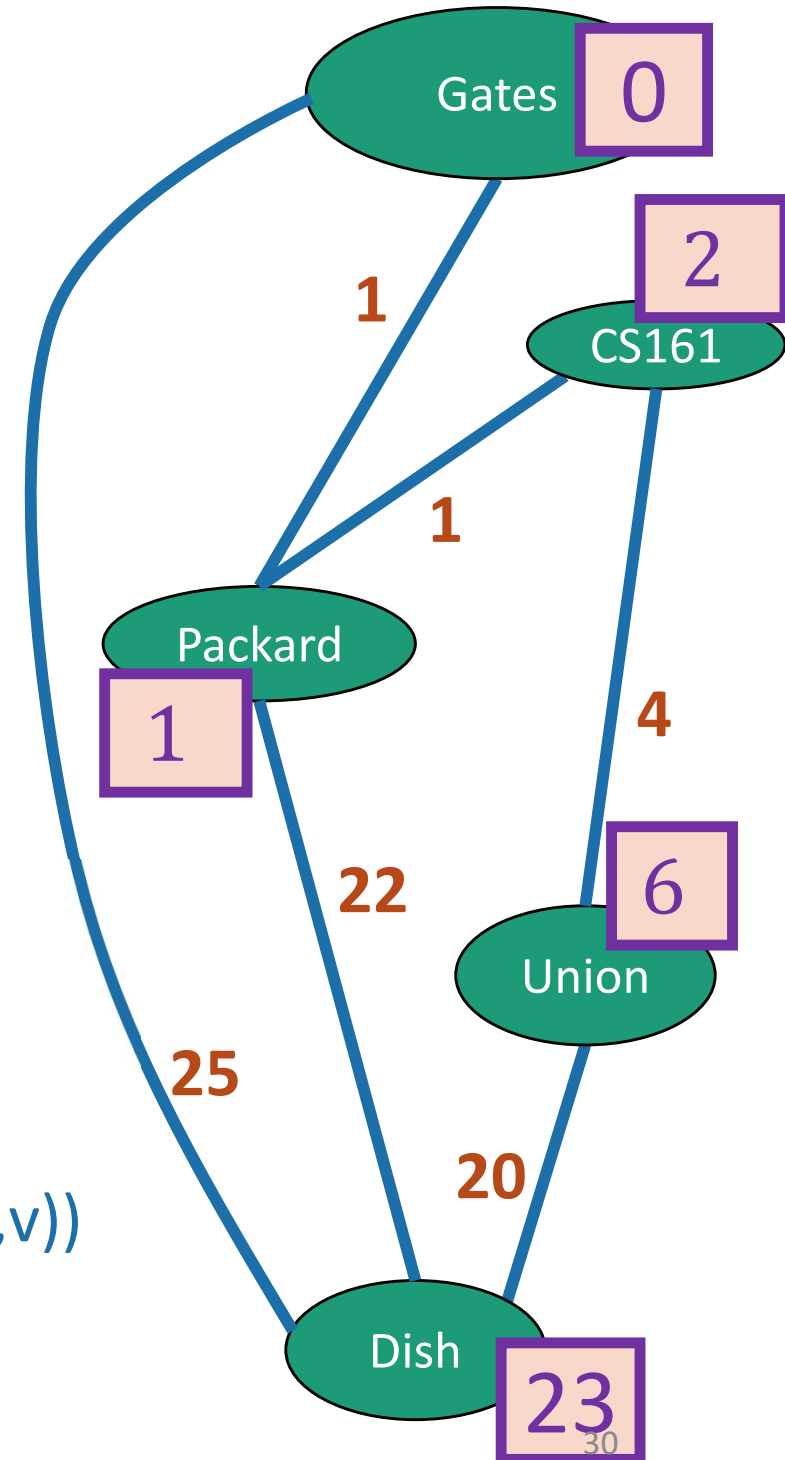


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.




Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat




Dijkstra by example

How far is a node from Gates?




I'm not sure yet



I'm sure

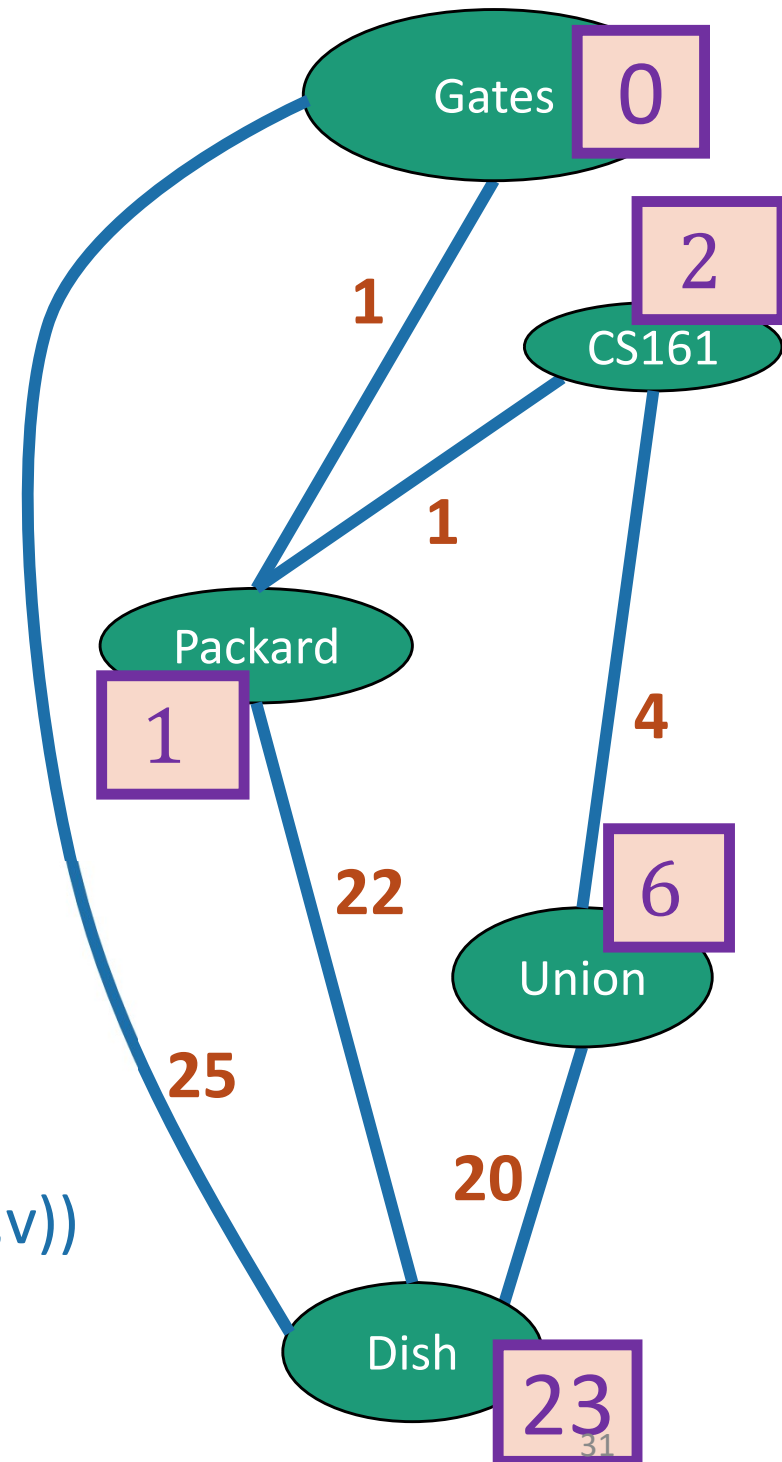


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.
- Repeat
- After all nodes are **sure**, say that $d(\text{Gates}, v) = d[v]$ for all v



Dijkstra's algorithm

Dijkstra(G,s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
 - **For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $d(s, v) = d[v]$

Why does this work?

- **Theorem:**

- Suppose we run Dijkstra on $G=(V,E)$, starting from s .
- At the end of the algorithm, the estimate $d[v]$ is the actual distance $d(s,v)$.

Let's rename "Gates" to "s", our starting vertex.

- **Proof outline:**

- **Claim 1:** For all v , $d[v] \geq d(s,v)$.
- **Claim 2:** When a vertex v is marked **sure**, $d[v] = d(s,v)$.

- **Claims 1 and 2 imply the theorem.**

- When v is marked **sure**, $d[v] = d(s,v)$.
- $d[v] \geq d(s,v)$ and never increases, so after v is **sure**, $d[v]$ stops changing.
- This implies that at any time *after* v is marked **sure**, $d[v] = d(s,v)$.
- All vertices are **sure** at the end, so all vertices end up with $d[v] = d(s,v)$.

Claim 2

Claim 1 + def of algorithm

Running time?

Dijkstra(G,s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $\text{dist}(s, v) = d[v]$

- n iterations (one per vertex)
- How long does one iteration take?

Depends on how we implement it...

We need a data structure that:

- Stores unsure vertices v
- Keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**.

Total running time is big-oh of:

$$\sum_{u \in V} \left(T(\text{findMin}) + \left(\sum_{v \in u.\text{neighbors}} T(\text{updateKey}) \right) + T(\text{removeMin}) \right)$$
$$= n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey})$$

If we use an array

- $T(\text{findMin}) = O(n)$
- $T(\text{removeMin}) = O(n)$
- $T(\text{updateKey}) = O(1)$
- Running time of Dijkstra
 - $= O(n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey}))$
 - $= O(n^2) + O(m)$
 - $= O(n^2)$


If we use a BST

- $T(\text{findMin}) = O(\log(n))$
- $T(\text{removeMin}) = O(\log(n))$
- $T(\text{updateKey}) = O(\log(n))$
- Running time of Dijkstra
 - $= O(n(T(\text{findMin}) + T(\text{removeMin}))) + m T(\text{updateKey})$
 - $= O(n \log(n)) + O(m \log(n))$
 - $= O((n + m) \log(n))$

Better than an array if the graph is sparse!
aka if m is much smaller than n^2

Heaps support these operations

- findMin
- removeMin
- updateKey




- A **heap** is a tree-based data structure that has the property that **every node has a smaller key than its children.**

Say we use a Heap

- $T(\text{findMin}) = O(1)$
- $T(\text{removeMin}) = O(\log(n))$
- $T(\text{updateKey}) = O(1)$
- Running time of Dijkstra
 - $= O(n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey}))$
 - $= O(n \log(n) + m)$

Ứng dụng thực tế

- eg, **OSPF (Open Shortest Path First)**, a routing protocol for IP networks, uses Dijkstra.



Hạn chế

- Chỉ áp dụng cho đồ thị có trọng số không âm (*non-negative edge weights*).
- Nếu một cạnh thay đổi trọng số, phải chạy lại thuật toán từ đầu.